

Experiment:3

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Steps:**Create 3 EC2 Amazon Linux Instances on AWS:**

- Created three instances (Master, worker-1, and worker-2) to set up the Kubernetes cluster.

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type
<input type="checkbox"/>	Master-Kuber	i-043b96c835c70d2d7	Running	t2.medium
<input type="checkbox"/>		i-0d46b0436cf21092b	Terminated	t2.micro
<input type="checkbox"/>	Worker-1-Kuber	i-0da6acf8052254ef8	Running	t2.medium
<input type="checkbox"/>	Worker-2-Kuber	i-0fbc853664be5368a	Running	t2.medium

- From now on, until mentioned, perform these steps on all 3 machines.
Install Docker on All Machines:

Commands:

```
sudo su
yum update -y
yum install docker -y
service docker start
systemctl enable docker
docker --version
```

Purpose: Docker is essential for running containers in Kubernetes. The installation commands ensure Docker is updated, installed, started, and set to start on boot.

```
[root@ip-172-31-19-73 home]# yum update -y
Last metadata expiration check: 0:02:58 ago on Sun Sep 15 04:37:38 2024.
Dependencies resolved.
Nothing to do.
Complete!
[root@ip-172-31-19-73 home]# yum install docker -y
Last metadata expiration check: 0:03:19 ago on Sun Sep 15 04:37:38 2024.
Package docker-25.0.6-1.amzn2023.0.2.x86_64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
[root@ip-172-31-19-73 home]# service docker start
Redirecting to /bin/systemctl start docker.service
[root@ip-172-31-19-73 home]# systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[root@ip-172-31-19-73 home]# docker --version
Docker version 25.0.5, build 5dc9bcc
[root@ip-172-31-19-73 home]# ||
```

i-043b96c835c70d2d7 (Master-Kuber)

```
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.
```

Verifying	:	containerd-1.7.20-1.amzn2023.0.1.x86_64	1/10
Verifying	:	docker-25.0.6-1.amzn2023.0.2.x86_64	2/10
Verifying	:	iptables-libs-1.8.8-3.amzn2023.0.2.x86_64	3/10
Verifying	:	iptables-nft-1.8.8-3.amzn2023.0.2.x86_64	4/10
Verifying	:	libcgroup-3.0-1.amzn2023.0.1.x86_64	5/10
Verifying	:	libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64	6/10
Verifying	:	libnftnl-1.0.1-19.amzn2023.0.2.x86_64	7/10
Verifying	:	libnftnl-1.2.2-2.amzn2023.0.2.x86_64	8/10
Verifying	:	pigz-2.5-1.amzn2023.0.3.x86_64	9/10
Verifying	:	runC-1.1.13-1.amzn2023.0.1.x86_64	10/10

Installed:

containerd-1.7.20-1.amzn2023.0.1.x86_64	docker-25.0.6-1.amzn2023.0.2.x86_64	iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
iptables-nft-1.8.8-3.amzn2023.0.2.x86_64	libcgroup-3.0-1.amzn2023.0.1.x86_64	libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
libnftnl-1.0.1-19.amzn2023.0.2.x86_64	libnftnl-1.2.2-2.amzn2023.0.2.x86_64	pigz-2.5-1.amzn2023.0.3.x86_64
runC-1.1.13-1.amzn2023.0.1.x86_64		

```
Complete!
[root@ip-172-31-29-38 ec2-user]# service docker start
Redirecting to /bin/systemctl start docker.service
[root@ip-172-31-29-38 ec2-user]# systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[root@ip-172-31-29-38 ec2-user]# docker --version
Docker version 25.0.5, build 5dc9bcc
[root@ip-172-31-29-38 ec2-user]#
```

i-0da6acf8052254ef8 (Worker-1-Kuber)

PublicIPs: 3.92.32.91 PrivateIPs: 172.31.29.38

```
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.
```

Verifying	:	containerd-1.7.20-1.amzn2023.0.1.x86_64	1/10
Verifying	:	docker-25.0.6-1.amzn2023.0.2.x86_64	2/10
Verifying	:	iptables-libs-1.8.8-3.amzn2023.0.2.x86_64	3/10
Verifying	:	iptables-nft-1.8.8-3.amzn2023.0.2.x86_64	4/10
Verifying	:	libcgroup-3.0-1.amzn2023.0.1.x86_64	5/10
Verifying	:	libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64	6/10
Verifying	:	libnftnl-1.0.1-19.amzn2023.0.2.x86_64	7/10
Verifying	:	libnftnl-1.2.2-2.amzn2023.0.2.x86_64	8/10
Verifying	:	pigz-2.5-1.amzn2023.0.3.x86_64	9/10
Verifying	:	runC-1.1.13-1.amzn2023.0.1.x86_64	10/10

Installed:

containerd-1.7.20-1.amzn2023.0.1.x86_64	docker-25.0.6-1.amzn2023.0.2.x86_64	iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
iptables-nft-1.8.8-3.amzn2023.0.2.x86_64	libcgroup-3.0-1.amzn2023.0.1.x86_64	libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
libnftnl-1.0.1-19.amzn2023.0.2.x86_64	libnftnl-1.2.2-2.amzn2023.0.2.x86_64	pigz-2.5-1.amzn2023.0.3.x86_64
runC-1.1.13-1.amzn2023.0.1.x86_64		

```
Complete!
[root@ip-172-31-28-21 ec2-user]# service docker start
Redirecting to /bin/systemctl start docker.service
[root@ip-172-31-28-21 ec2-user]# systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[root@ip-172-31-28-21 ec2-user]# docker --version
Docker version 25.0.5, build 5dc9bcc
[root@ip-172-31-28-21 ec2-user]#
```

i-0fbc853664be5368a (Worker-2-Kuber)

PublicIPs: 50.19.34.228 PrivateIPs: 172.31.28.21

Then, configure cgroup in a daemon.json file.

```
cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

Purpose: Configures Docker to use the systemd cgroup driver, which is recommended for Kubernetes, and ensures Docker logs are managed and storage is optimized.

```
"log-driver": "json-file",
"log-opts": {
  "max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
[root@ip-172-31-19-73 docker]# cat <<EOF > /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
> EOF
[root@ip-172-31-19-73 docker]# sudo $?
sudo: 0: command not found
[root@ip-172-31-19-73 docker]# systemctl enable docker
[root@ip-172-31-19-73 docker]# systemctl daemon-reload
[root@ip-172-31-19-73 docker]# systemctl restart docker
[root@ip-172-31-19-73 docker]# docker info | grep -i cgroup
Cgroup Driver: systemd
Cgroup Version: 2
cgroupns
[root@ip-172-31-19-73 docker]#
```

i-043b96c835c70d2d7 (Master-Kuber)

PublicIPs: 54.242.131.150 PrivateIPs: 172.31.19.73

```
[root@ip-172-31-29-38 ec2-user]# cd /etc/docker
[root@ip-172-31-29-38 docker]# cat <<EOF > /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
[root@ip-172-31-29-38 docker]# systemctl enable docker
[root@ip-172-31-29-38 docker]# systemctl daemon-reload
[root@ip-172-31-29-38 docker]# systemctl restart docker
[root@ip-172-31-29-38 docker]# docker info | grep -i cgroup
Cgroup Driver: systemd
Cgroup Version: 2
cgroupns
[root@ip-172-31-29-38 docker]#
```

i-Oda6acf8052254ef8 (Worker-1-Kuber)

PublicIPs: 3.92.32.91 PrivateIPs: 172.31.29.38

```
[root@ip-172-31-28-21 ec2-user]# cd /etc/docker
[root@ip-172-31-28-21 docker]# cat <<EOF > /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
[root@ip-172-31-28-21 docker]# systemctl enable docker
[root@ip-172-31-28-21 docker]# systemctl daemon-reload
[root@ip-172-31-28-21 docker]# systemctl restart docker
[root@ip-172-31-28-21 docker]# docker info | grep -i cgroup
Cgroup Driver: systemd
Cgroup Version: 2
cgroupns
[root@ip-172-31-28-21 docker]#
```

i-Ofbc853664be5368a (Worker-2-Kuber)

PublicIPs: 50.19.34.228 PrivateIPs: 172.31.28.21

1. Set SELinux to permissive mode:
These instructions are for Kubernetes 1.31.

```
sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```
2. Add the Kubernetes yum repository. The exclude parameter in the repository definition ensures that the packages related to Kubernetes are not upgraded upon running yum update as there's a special procedure that must be followed for upgrading Kubernetes. Please note that this repository have packages only for Kubernetes 1.31; for other Kubernetes minor versions, you need to change the Kubernetes minor version in the URL to match your desired minor version (you should also check that you are reading the documentation for the version of Kubernetes that you plan to install).

```
cat <<EOF | sudo tee /etc/yum/repos.d/kubernetes.repo
[kubernetes]
```

```
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

3. Install kubelet, kubeadm and kubectl:
 sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
4. (Optional) Enable the kubelet service before running kubeadm:
 sudo systemctl enable --now kubelet

```
[root@ip-172-31-19-73 docker]# setenforce 0
sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[root@ip-172-31-19-73 docker]# cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[root@ip-172-31-19-73 docker]# yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Kubernetes
Package kubelet-1.31.1-150500.1.1.x86_64 is already installed.
Package kubeadm-1.31.1-150500.1.1.x86_64 is already installed.
Package kubectl-1.31.1-150500.1.1.x86_64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
[root@ip-172-31-19-73 docker]# systemctl enable --now kubelet
[root@ip-172-31-19-73 docker]# ||
```

i-043b96c835c70d2d7 (Master-Kuber)

PublicIPs: 54.242.131.150 PrivateIPs: 172.31.19.73


```

Installing      : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
Installing      : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
Installing      : kubelet-1.31.1-150500.1.1.x86_64
Running scriptlet: kubelet-1.31.1-150500.1.1.x86_64
Installing      : kubeadm-1.31.1-150500.1.1.x86_64
Installing      : kubectrl-1.31.1-150500.1.1.x86_64
Running scriptlet: kubectrl-1.31.1-150500.1.1.x86_64
Verifying       : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
Verifying       : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
Verifying       : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
Verifying       : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64
Verifying       : cri-tools-1.31.1-150500.1.1.x86_64
Verifying       : kubeadm-1.31.1-150500.1.1.x86_64
Verifying       : kubectrl-1.31.1-150500.1.1.x86_64
Verifying       : kubelet-1.31.1-150500.1.1.x86_64
Verifying       : kubernetes-cni-1.5.1-150500.1.1.x86_64

Installed:
  conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.31.1-150500.1.1.x86_64
  kubectrl-1.31.1-150500.1.1.x86_64               kubelet-1.31.1-150500.1.1.x86_64
  libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64  libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64

Complete!
[root@ip-172-31-29-38 docker]# systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service
[root@ip-172-31-29-38 docker]# | |

```

i-0da6acf8052254ef8 (Worker-1-Kuber)

PublicIPs: 3.92.32.91 PrivateIPs: 172.31.29.38

```

Installing      : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
Installing      : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
Installing      : kubelet-1.31.1-150500.1.1.x86_64
Running scriptlet: kubelet-1.31.1-150500.1.1.x86_64
Installing      : kubeadm-1.31.1-150500.1.1.x86_64
Installing      : kubectrl-1.31.1-150500.1.1.x86_64
Running scriptlet: kubectrl-1.31.1-150500.1.1.x86_64
Verifying       : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
Verifying       : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
Verifying       : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
Verifying       : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64
Verifying       : cri-tools-1.31.1-150500.1.1.x86_64
Verifying       : kubeadm-1.31.1-150500.1.1.x86_64
Verifying       : kubectrl-1.31.1-150500.1.1.x86_64
Verifying       : kubelet-1.31.1-150500.1.1.x86_64
Verifying       : kubernetes-cni-1.5.1-150500.1.1.x86_64

Installed:
  conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.31.1-150500.1.1.x86_64
  kubectrl-1.31.1-150500.1.1.x86_64               kubelet-1.31.1-150500.1.1.x86_64
  libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64  libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64

Complete!
[root@ip-172-31-28-21 docker]# systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service
[root@ip-172-31-28-21 docker]# | |

```

i-0fbcb853664be5368a (Worker-2-Kuber)

PublicIPs: 50.19.34.228 PrivateIPs: 172.31.28.21

5. After installing Kubernetes, we need to configure internet options to allow bridging.

```
sudo swapoff -a
```

```
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
```

```
sudo sysctl -p
```

Purpose: Disables swap memory and configures network settings to ensure proper networking and performance for Kubernetes.

```
[ec2-user@ip-172-31-19-73 ~]$ swapoff -a
[ec2-user@ip-172-31-19-73 ~]$ echo "net.bridge
-bash: /etc/sysctl.conf: Permission denied
[ec2-user@ip-172-31-19-73 ~]$ echo "net.bridge
net.bridge.bridge-nf-call-iptables=1
[ec2-user@ip-172-31-19-73 ~]$ sudo sysctl -p
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-19-73 ~]$
```

i-043b96c835c70d2d7 (Master-Kuber)

PublicIPs: 54.242.131.150 PrivateIPs: 172.31.19.73

```
[root@ip-172-31-29-38 docker]# swapoff -a
[root@ip-172-31-29-38 docker]# echo "net.bri
net.bridge.bridge-nf-call-iptables=1
[root@ip-172-31-29-38 docker]# sudo sysctl -
net.bridge.bridge-nf-call-iptables = 1
[root@ip-172-31-29-38 docker]#
```

i-0da6acf8052254ef8 (Worker-1-Kuber)

PublicIPs: 3.92.32.91 PrivateIPs: 172.31.29.38

```
[root@ip-172-31-28-21 docker]# swapoff -a
[root@ip-172-31-28-21 docker]# echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
net.bridge.bridge-nf-call-iptables=1
[root@ip-172-31-28-21 docker]# sudo sysctl -p
net.bridge.bridge-nf-call-iptables = 1
[root@ip-172-31-28-21 docker]#
```

6. Perform this ONLY on the Master machine

kubeadm init first runs a series of prechecks to ensure that the machine is ready to run Kubernetes. These prechecks expose warnings and exit on errors. kubeadm init then downloads and installs the cluster control plane components. This may take several minutes. After it finishes you should see:

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a Pod network to the cluster.

Run

"kubectl apply -f

<https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml>"

with one of the options listed at : [/docs/concepts/cluster-administration/addons/](#)

Purpose: Initialises the Kubernetes master, sets up the kubeconfig file for cluster access, and deploys the Flannel network plugin for pod networking.

You can now join any number of machines by running the following on each node as root:

```
[root@ip-172-31-19-73 /]# mkdir -p $HOME/.kube
[root@ip-172-31-19-73 /]# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
cp: overwrite '/root/.kube/config'? y
[root@ip-172-31-19-73 /]# sudo chown $(id -u):$(id -g) $HOME/.kube/config
[root@ip-172-31-19-73 /]# kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
namespace/kube-flannel created
serviceaccount/flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
[root@ip-172-31-19-73 /]# kubectl get pods
No resources found in default namespace.
[root@ip-172-31-19-73 /]#
[root@ip-172-31-19-73 /]# kubectl get pods --all-namespaces
NAMESPACE      NAME                                                    READY   STATUS              RESTARTS   AGE
kube-flannel    kube-flannel-ds-lbgbd                                  0/1     CrashLoopBackOff    3 (44s ago) 99s
kube-system     coredns-7c65d6cfc9-8vjkk                              0/1     ContainerCreating    0           3m13s
kube-system     coredns-7c65d6cfc9-bl78b                              0/1     ContainerCreating    0           3m13s
kube-system     etcd-ip-172-31-19-73.ec2.internal                     1/1     Running             19 (3m5s ago) 3m21s
kube-system     kube-apiserver-ip-172-31-19-73.ec2.internal            1/1     Running             20 (3m25s ago) 3m21s
kube-system     kube-controller-manager-ip-172-31-19-73.ec2.internal   1/1     Running             21 (3m5s ago) 3m21s
kube-system     kube-proxy-j9fdq                                       0/1     CrashLoopBackOff    2 (16s ago)  3m13s
kube-system     kube-scheduler-ip-172-31-19-73.ec2.internal            1/1     Running             21 (3m ago)   3m57s
[root@ip-172-31-19-73 /]# kubectl get nodes
NAME                                STATUS    ROLES    AGE     VERSION
ip-172-31-19-73.ec2.internal        Ready    control-plane  4m30s  v1.31.1
```

sudo yum install iproute-tc socat -y

sudo systemctl enable kubelet

sudo systemctl restart kubelet

Purpose: Installs necessary packages on worker nodes, ensures **kubelet** starts, and joins the nodes to the Kubernetes cluster using the provided token and hash.

```
[root@ip-172-31-29-38 docker]# sudo yum install iproute-tc socat -y
sudo systemctl enable kubelet
sudo systemctl restart kubelet
Last metadata expiration check: 0:57:00 ago on Sun Sep 15 04:57:12 2024.
Dependencies resolved.

Package                               Architecture Version
Installing:
iproute-tc                           x86_64    5.10.0-2.amzn
socat                                 x86_64    1.7.4.2-1.amzn

Transaction Summary
--
Install 2 Packages

Total download size: 758 k
Installed size: 2.0 M
Downloading Packages:
(1/2): iproute-tc-5.10.0-2.amzn2023.0.5.x86_64.rpm
(2/2): socat-1.7.4.2-1.amzn2023.0.2.x86_64.rpm

Total
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
```


Paste this to both worker node to join:

```
kubeadm join 172.31.19.73:6443 --token cenfmg.0y3m7m51649vj3bd \
--discovery-token-ca-cert-hash
sha256:3443d69f393c593e9279c6e26ea7b37208db3db34bb1507c01f790a57f4b8a7b
```

```
[root@ip-172-31-19-73 /]# kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-19-73.ec2.internal        Ready    control-plane   11m   v1.31.1
[root@ip-172-31-19-73 /]#
```

i-043b96c835c70d2d7 (Master-Kuber)

PublicIPs: 54.242.131.150 PrivateIPs: 172.31.19.73

```
[root@ip-172-31-29-38 docker]# kubeadm join 172.31.19.73:6443 --token cenfmg.0y3m7m51649vj3bd \
--discovery-token-ca-cert-hash sha256:3443d69f393c593e9279c6e26ea7b37208db3db34bb1507c01f790a57f4b8a7b
[preflight] Running pre-flight checks
^C
[root@ip-172-31-29-38 docker]#
```

i-0da6acf8052254ef8 (Worker-1-Kuber)

PublicIPs: 3.92.32.91 PrivateIPs: 172.31.29.38

```
[root@ip-172-31-28-21 docker]# kubeadm join 172.31.19.73:6443 --token cenfmg.0y3m7m51649vj3bd \
--discovery-token-ca-cert-hash sha256:3443d69f393c593e9279c6e26ea7b37208db3db34bb1507c01f790a57f4b8a7b
[preflight] Running pre-flight checks
^C
[root@ip-172-31-28-21 docker]#
```

i-0fbc853664be5368a (Worker-2-Kuber)

PublicIPs: 50.19.34.228 PrivateIPs: 172.31.28.21

Conclusion :

In this experiment, we aimed to understand Kubernetes cluster architecture by setting up a Kubernetes cluster on AWS EC2 instances. We started by creating three Amazon Linux instances (Master, worker-1, and worker-2). Docker was installed and configured on all nodes, SELinux was set to permissive mode, and the Kubernetes repository was added. We then installed the Kubernetes packages and initialised the cluster on the master node. After configuring networking and deploying the Flannel plugin, we joined the worker nodes to the cluster. Despite following the steps, the setup did not complete successfully, indicating issues with establishing a fully functional Kubernetes cluster with all nodes connected.